Five Proposals for Deployed Compute

Sustainable Market Mechanisms for Information Streams and Aspirational Guardrail Principles

Nicholas J. Teague, PE (Orlando, FL)

Drafted with the assistance of Grok (a LLM from xAI). The human author retains responsibility for the content.

Abstract

These proposals outline a layered market framework for allocating access to high-value information streams, where commercial deployment of potentially fault-tolerant systems risk undermining mission-critical integrity. Identifiable streams are the primary scarce good. Access is organized in tiers: Unrestricted raw-feed-to-raw-compute paths remain reserved for the highest tiers. Supervisory systems rely on auctions of nested exclusive participants retaining economic competition among high-stakes participants and enabling adaptive conventions towards downstream use, a middle tier of qualified sampling-rate auctions with sample rate distribution envelopes (where sampling rate deviations may serve as a “pressure relief valve” to mitigate emergent fragile macro states), and lower tiers governed by raw system availability or explicit acceptance of rate and availability deviations (suitable for experimental or academic use). Private deviation-weighted metrics supplement binary availability. The design enables commercial demand revelation while preserving isolation and qualification for high-consequence systems. Directions for further extensions may enable compute marketplaces, framing future aspirational principles for the field. The overall aim is to support evaluation of whether these conventions can help protect mission-critical integrity while still enabling commercial and experimental demand revelation, and to extend the thinking toward more sustainable forms of competition among sources of scaled compute.

“The wind does not break a tree that bends.” — African proverb

---

These five deployment proposals are offered to support evaluation of merits. They remain speculative on implementation detail and do not constitute a finished policy or call for immediate adoption.

Point 1 — Foundational Conventions for Information-Stream Access

The core problem is that commercial deployment of potentially fault-tolerant systems, absent disciplined prioritization among use cases, can jeopardize the integrity of mission-critical systems through shared components, standards, talent, and data paths. The primary scarce and excludable good is therefore the right to sample identifiable information streams (satellite feeds, sensor outputs, telemetry), which are more realistic to gate than fungible raw compute. Access is rationed by qualified auctions whose eligibility favors participants who already aggregate operational context and stakes, so that prioritization reflects informed assessments of importance rather than pure economic wager. AI agents make low-latency, multi-attribute bidding practical. A supervisory compute layer holds higher default visibility and reserved rights insulated from commercial manipulation. A binding industry convention is that unrestricted raw-feed-to-raw-compute paths remain reserved for the highest-tier mission-critical settings only; commercial users consume rate-limited or derived products. Lightweight settlement suffices for commercial transfers; heavier integrity machinery is reserved for the supervisory layer. These conventions create a market signal from revealed demand back to sensor deployment while preserving integrity for high-consequence uses.

Point 2 — Sampling-Rate Envelopes as Pressure-Relief Mechanisms

Sampling-rate auctions may clear not only fixed nominal rates but nominal rates within guaranteed envelopes (for example, a design rate expandable or contractible within published bounds for bounded intervals under supervisory invocation, with defined recovery). This may function as a pressure-relief valve analogous to the controlled voltage and power-quality bands used in electricity transmission systems, and thus serve to mitigate emergent fragile macro states. Commercial bidders price both the expected nominal performance and the statistical properties of the envelope. Supervisory systems can thereby accommodate transient integrity needs without hard pre-emption of commercial clearings. The envelope must remain narrow enough, and its exercise visible and accountable enough, that commercial confidence in the

Copyright © Nicholas Teague 2026

product is preserved. All foundational elements of Point 1—stream-level gating, qualification, agent mediation, supervisory reserved rights, and the raw-feed restriction—continue to apply.

Point 3 — Deviation-Weighted Metrics and Intermediate Availability Signals

Binary availability (simple uptime ratio) is insufficient once rate deviations inside an envelope are permitted. A private secondary metric—deviation-weighted availability—penalizes larger departures from the design rate more heavily than mild ones. A natural form, resembling capacity-weighted availability, voltage-penalty functions, and integral-of-squared-error criteria, is $A_w = 1 - (1/T) \int w(\delta(t))\, dt$, where $\delta(t)$ is relative deviation and $w(\delta)$ rises quadratically or in tiers. Qualified users compute this (or any alternative fidelity measure) with their own methods or loss functions; the metric need not be encoded in auction rules. It informs their bids and allows ranking of systems by what stability is implied by their deviation behavior. In the present environment, where no common expectation of universal commercial uptime exists, ordinary material deviations from stated availability already serve as an intermediate, observable signal that tighter envelopes or tier separations may be warranted for high-consequence uses while fuller mechanisms are evaluated. This allows a graduated path to integrity preservation without abrupt connection constraints.

Point 4 — Multi-Tier Access Privileges

Access privileges are organized in explicit tiers. The higher layer consists of nested exclusive auctions on the supervisory path, limited to tightly qualified participants who carry operational responsibility or residual liability. These auctions retain economic competition and price discovery inside mission-critical settings and provide a venue for evaluating and adapting conventions (envelope parameters, weighting functions, qualification criteria) through pilot or shadow clearings. The middle tier comprises the qualified sampling-rate auctions with optional priced envelopes described above. The lower layer offers paths governed by raw system availability or by explicit acceptance of both sampling-rate and availability deviations; these paths are subservient to higher-layer allocations and are suited to academic or experimental settings in which conventions remain exploratory. Unrestricted raw-feed-to-raw-compute remains reserved for the highest tiers. Private deviation-weighted metrics and ordinary availability deviations remain as evaluation and signaling tools across tiers. The multi-tier structure keeps commercial and experimental demand revelation intact while ensuring that higher-consequence uses retain both priority and adaptive capacity.

Point 5 — Aspirational Principles for Concurrent Compute Marketplaces with Generative AI Guardrails

Looking further ahead, the author sees potential value in concurrent marketplaces among sources of compute, each with access to different sources of context, stakes assessments, and user profiles distributed across clouds. Such arrangements could support a more sustainable form of economic competition between compute vendors than token generation or raw FLOPs accessibility alone. The marketplace itself need not have defined boundaries in this proposal; the recognition is simply that future guardrails for generative AI will, in the author's opinion, require access to sources of context, stakes assessment, and user profiles external to any specific prompt window to enable effective guardrails.

In this connection the author has recommended, in the ChatFlavors project, that the types of guardrails implementable in this fashion could be constrained to a simple set of objectives designed to ensure alignment to the benefit of the users and their surroundings. ChatFlavors provides a template of default prompt injections defining a small number of fundamental interaction flavors (Biz, Fun, Res, Art, Cod, Sys, Cch, Tut, Doc), each carrying domain-appropriate nudges that grow stronger with higher stakes, while deferring to human decisions where humans are aware of consequences. The project's underlying philosophy holds that prompts alone will never suffice for workable guardrails and that consumer portals can usefully compete in their backends on how they evaluate stakes, user profiles, and context. The repository is available at https://github.com/Automunge/ChatFlavors.

These principles are offered as a potential direction for the field to iterate toward. Realization will require further infrastructure development and collaborative standardization, yet the core aspiration remains that compute competition can incorporate the quality and accessibility of external context, user profiles, and stake information without abandoning the commodity character of scaling processing capacity itself.

Copyright © Nicholas Teague 2026